

# Neural networks for text

# What is text?

**You can think of text as a sequence of**

- Characters
- **Words**
- Phrases and named entities
- Sentences
- Paragraphs
- ...

# Bag of words way (sparse)

~100k columns

	good	movie	very	a	did	like
very →	0	0	1	0	0	0
good →	1	0	0	0	0	0
movie →	0	1	0	0	0	0


# Bag of words way (sparse)

~100k columns

	good	movie	very	a	did	like
very	0	0	1	0	0	0
good	1	0	0	0	0	0
movie	0	1	0	0	0	0
very good movie	1	1	1	0	0	0

Bag of words representation  
is a sum of sparse one-hot-encoded vectors

# Neural way (dense)




## Word2vec property:

Words that have similar context tend to have collinear vectors


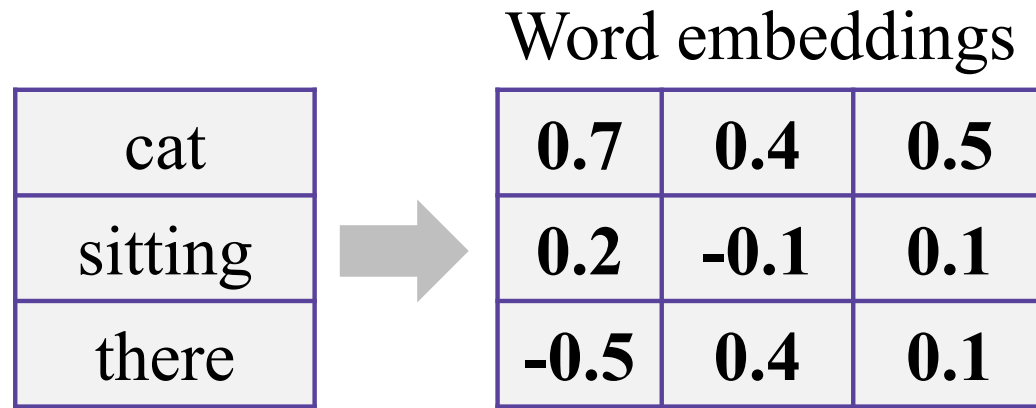
# Neural way (dense)

~300 columns




Sum of word2vec vectors  
can be a good text descriptor already!

# A better way: 1D convolutions



How do we make 2-grams?

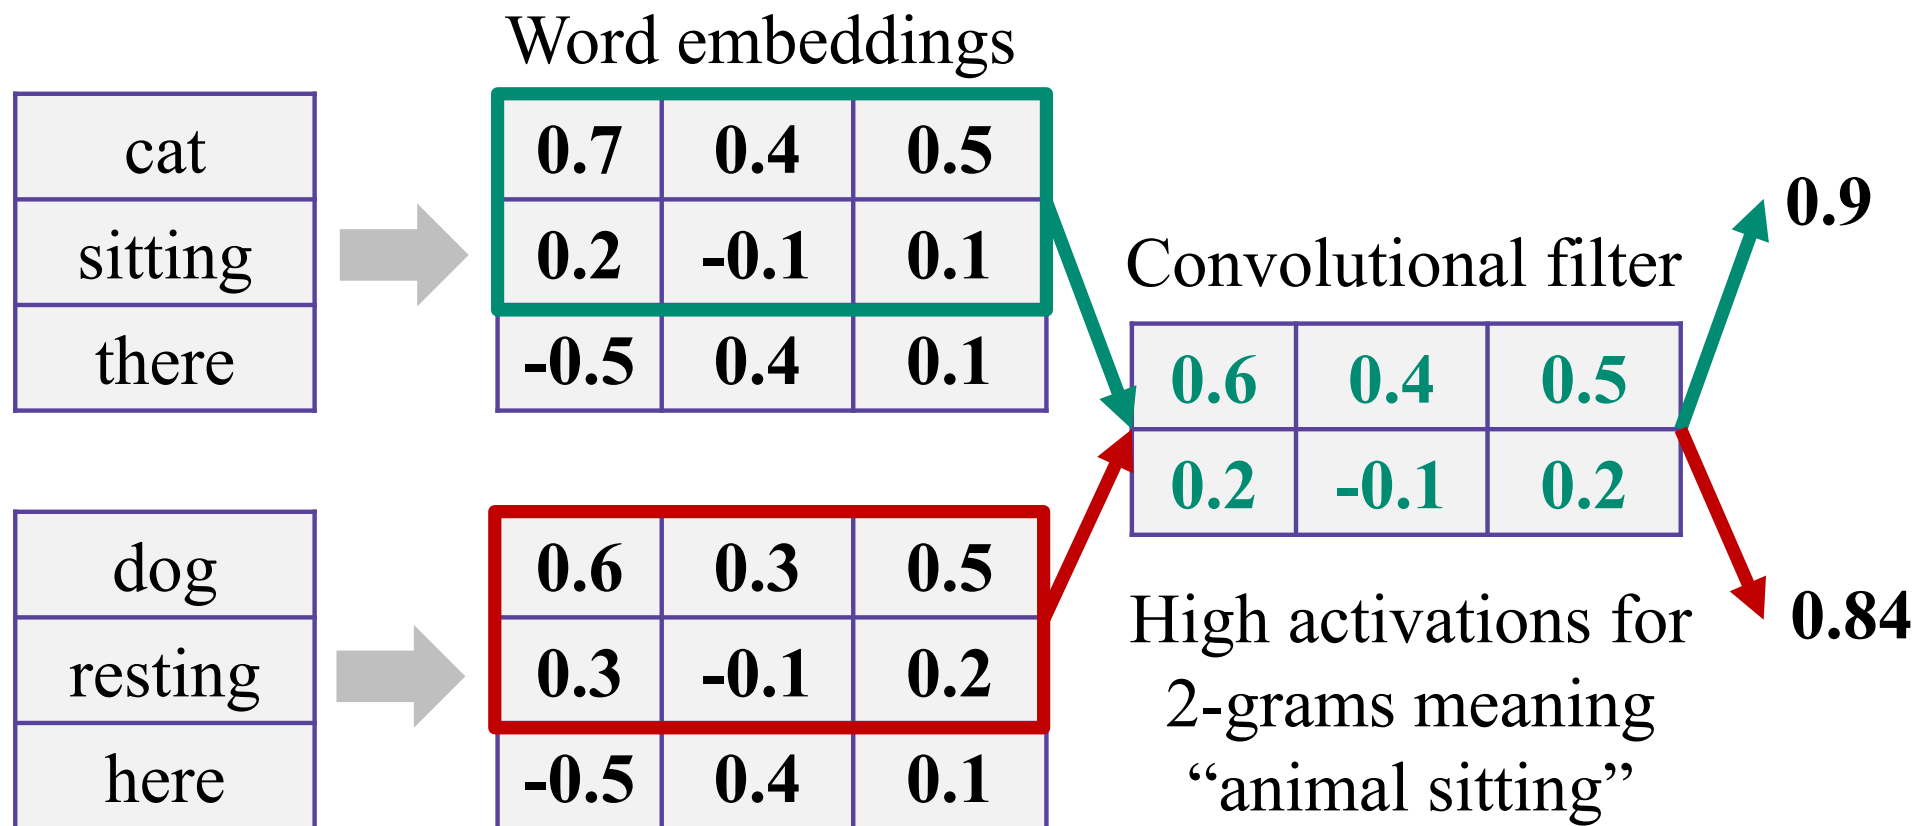
# A better way: 1D convolutions



- This convolution provides high activations for 2-grams with certain meaning



# A better way: 1D convolutions



- This convolution provides high activations for 2-grams with certain meaning
- Word2vec vectors for similar words are similar in terms of cosine distance (similar to dot product)


# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction




# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction



# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction




# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction




# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction



# 1D convolutions

- Can be extended to 3-grams, 4-grams, etc.
- One filter is not enough, need to track many n-grams
- They are called 1D because we slide the window only in one direction




# Let's train many filters

## Final architecture

- 3,4,5-gram windows with 100 filters each
- MLP on top of these 300 features

## Quality comparison on customer reviews (CR)

- Naïve Bayes on top of 1,2-grams – 86.3% accuracy
- 1D convolutions with MLP – 89.6% (+3.8%) accuracy





**Going deeper with text**

# What is text?

**You can think of text as a sequence of**

- **Characters**
- Words
- Phrases and named entities
- Sentences
- Paragraphs
- ...

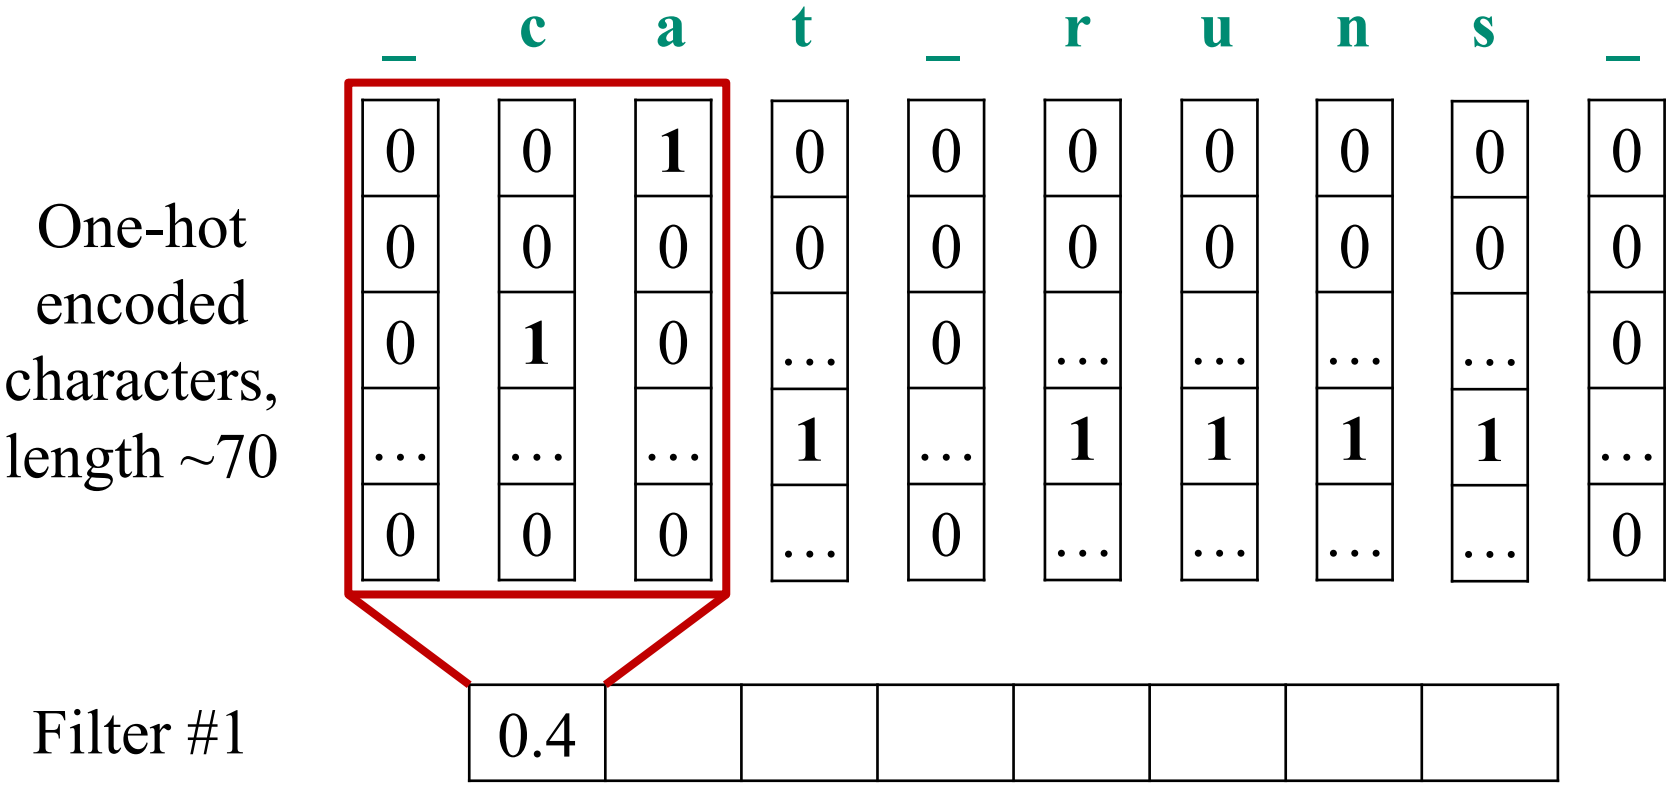
# Text as a sequence of characters

One-hot  
encoded  
characters,  
length  $\sim 70$

	<b>_</b>	<b>c</b>	<b>a</b>	<b>t</b>	<b>_</b>	<b>r</b>	<b>u</b>	<b>n</b>	<b>s</b>	<b>_</b>
	0	0	1	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0
	0	1	0	...	0	...	...	...	...	0
	...	...	...	1	...	1	1	1	1	...
	0	0	0	...	0	...	...	...	...	0

Let's start with character  $n$ -grams!

# 1D convolutions on characters



# 1D convolutions on characters

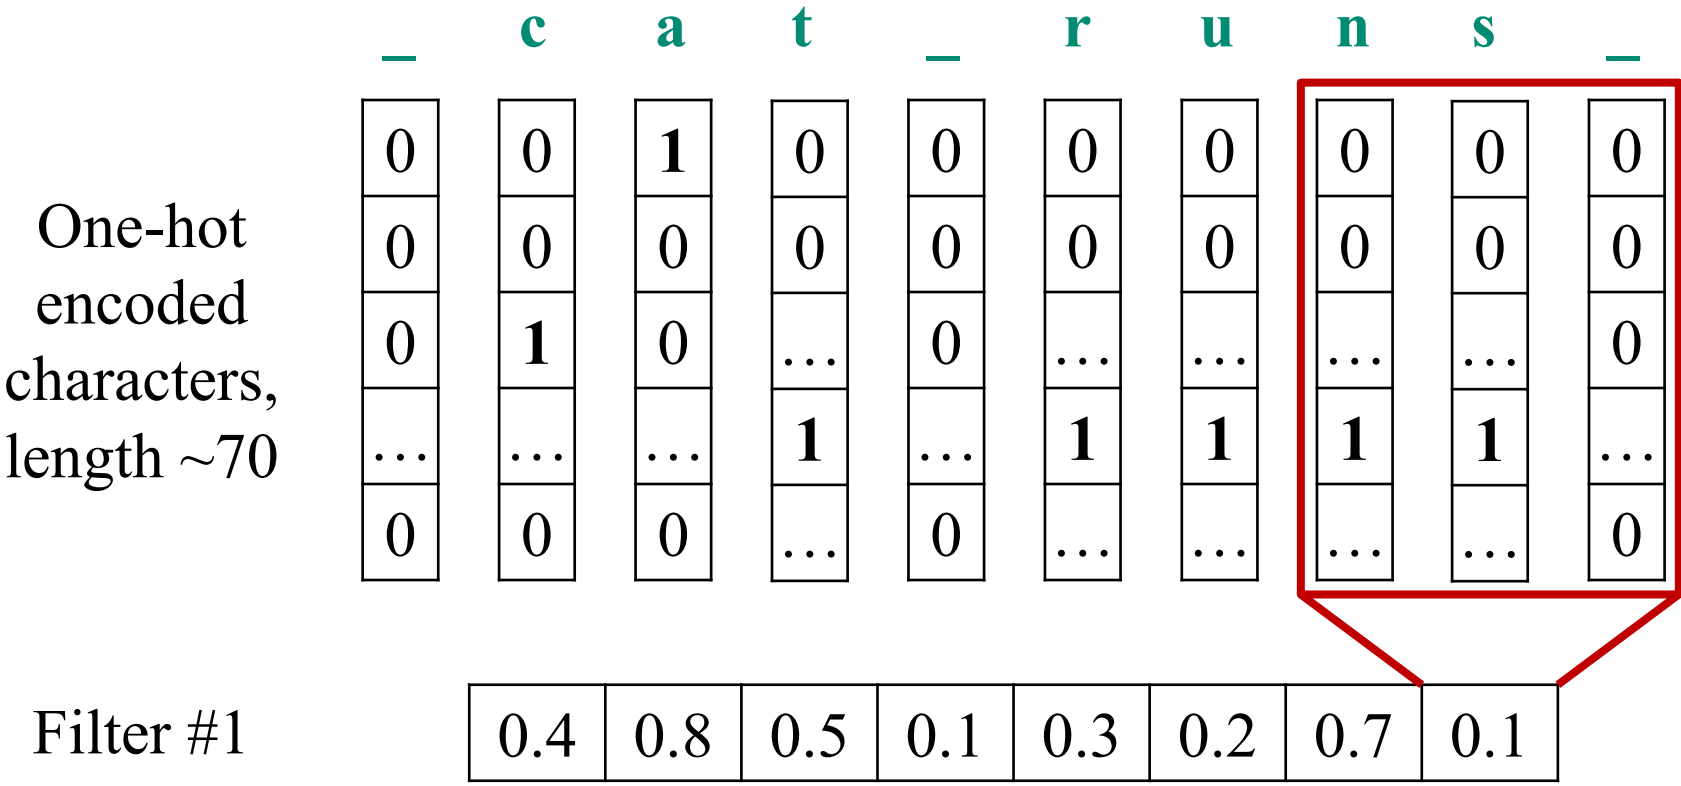
One-hot  
encoded  
characters,  
length ~70

_	c	a	t	_	r	u	n	s	_
0	0	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	1	0	...	0	...	...	...	...	0
...	...	...	1	...	1	1	1	1	...
0	0	0	...	0	...	...	...	...	0

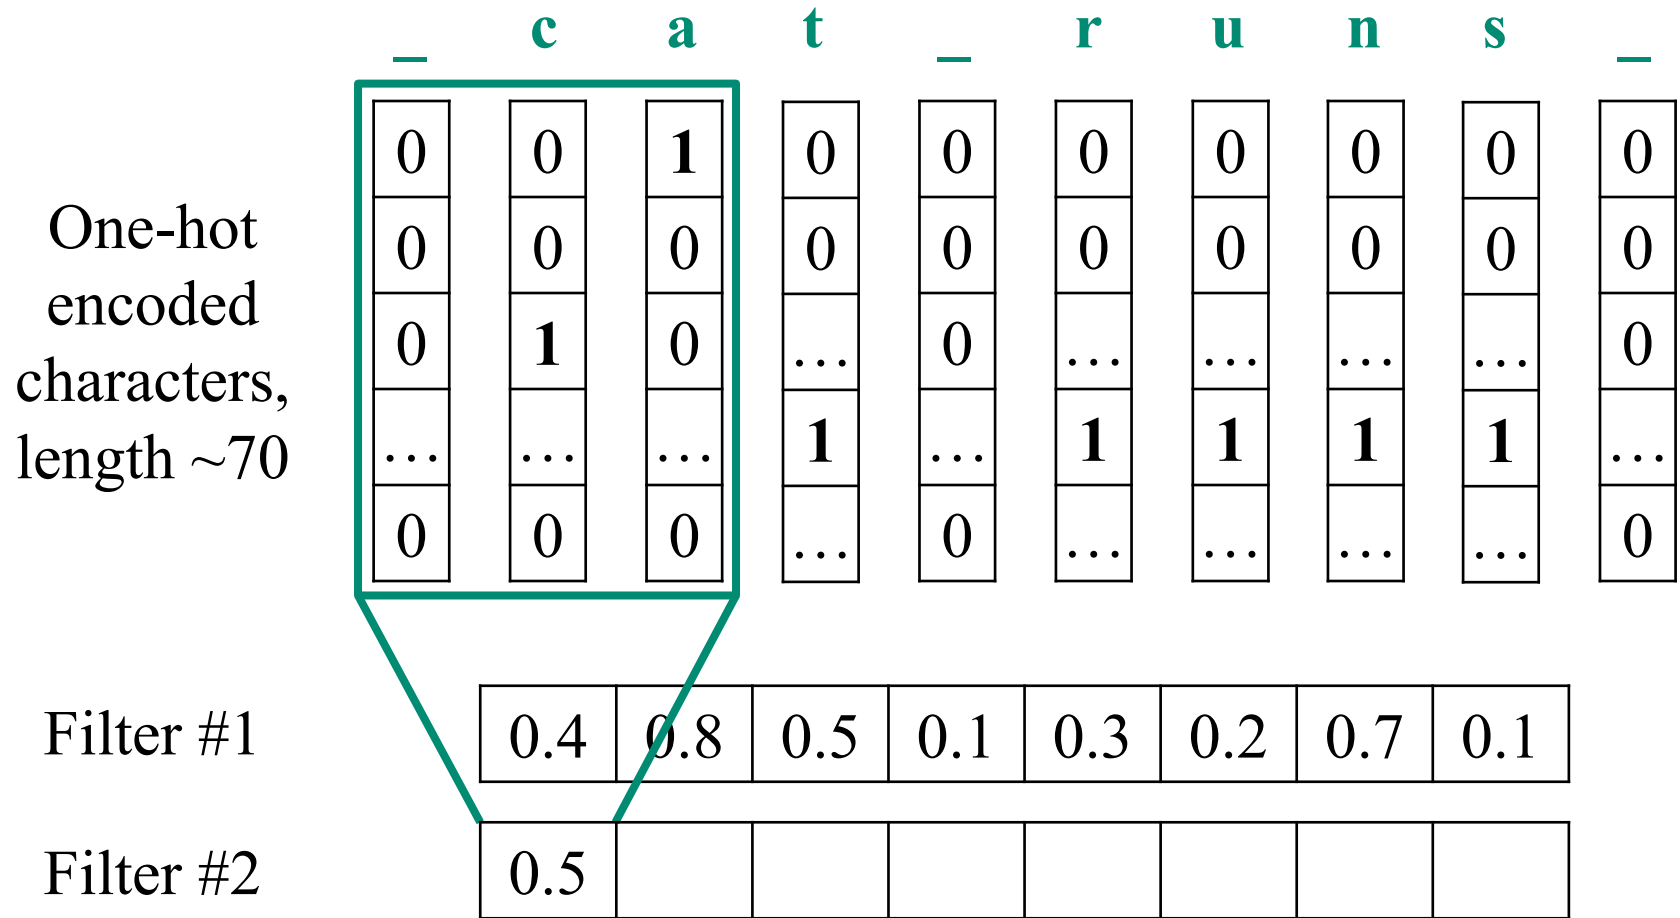
Filter #1

0.4	0.8						
-----	-----	--	--	--	--	--	--


# 1D convolutions on characters



# 1D convolutions on characters



# 1D convolutions on characters





# 1D convolutions on characters

One-hot encoded characters, length ~70	_	c	a	t	_	r	u	n	s	_
	0	0	1	0	0	0	0	0	0	0
	0	0	0	0	0	0	0	0	0	0
	0	1	0	...	0	...	...	...	...	0
	...	...	...	1	...	1	1	1	1	...
	0	0	0	...	0	...	...	...	...	0

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

~1024 filters

What's next? Let's add pooling!

# Max pooling

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

0.8			
-----	--	--	--

--	--	--	--

--	--	--	--

# Max pooling

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

0.8	0.5		
-----	-----	--	--

--	--	--	--

--	--	--	--

# Max pooling

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

0.8	0.5	0.3	0.7
-----	-----	-----	-----

--	--	--	--

--	--	--	--

# Max pooling

Filter #1

0.4	0.8	0.5	0.1	0.3	0.2	0.7	0.1
-----	-----	-----	-----	-----	-----	-----	-----

Filter #2

0.5	0.1	0.4	0.2	0.3	0.9	0.1	0.8
-----	-----	-----	-----	-----	-----	-----	-----

Filter #3

0.4	0.7	0.3	0.7	0.5	0.5	0.9	0.4
-----	-----	-----	-----	-----	-----	-----	-----

Pooling  
output

0.8	0.5	0.3	0.7
-----	-----	-----	-----

0.5	0.4	0.9	0.8
-----	-----	-----	-----

0.7	0.7	0.5	0.9
-----	-----	-----	-----

Provides a little bit of position  
invariance for character n-grams

# Repeat 1D convolution + pooling

Pooling  
output

0.8	0.5	0.3	0.7
0.5	0.4	0.9	0.8
0.7	0.7	0.5	0.9

Another filter #1

0.4	0.8	0.4	0.9
-----	-----	-----	-----

Another filter #2

0.9	0.8	0.6	0.5
-----	-----	-----	-----

# Repeat 1D convolution + pooling

Pooling  
output

0.8	0.5	0.3	0.7
0.5	0.4	0.9	0.8
0.7	0.7	0.5	0.9

Another filter #1

0.4	0.8	0.4	0.9
-----	-----	-----	-----

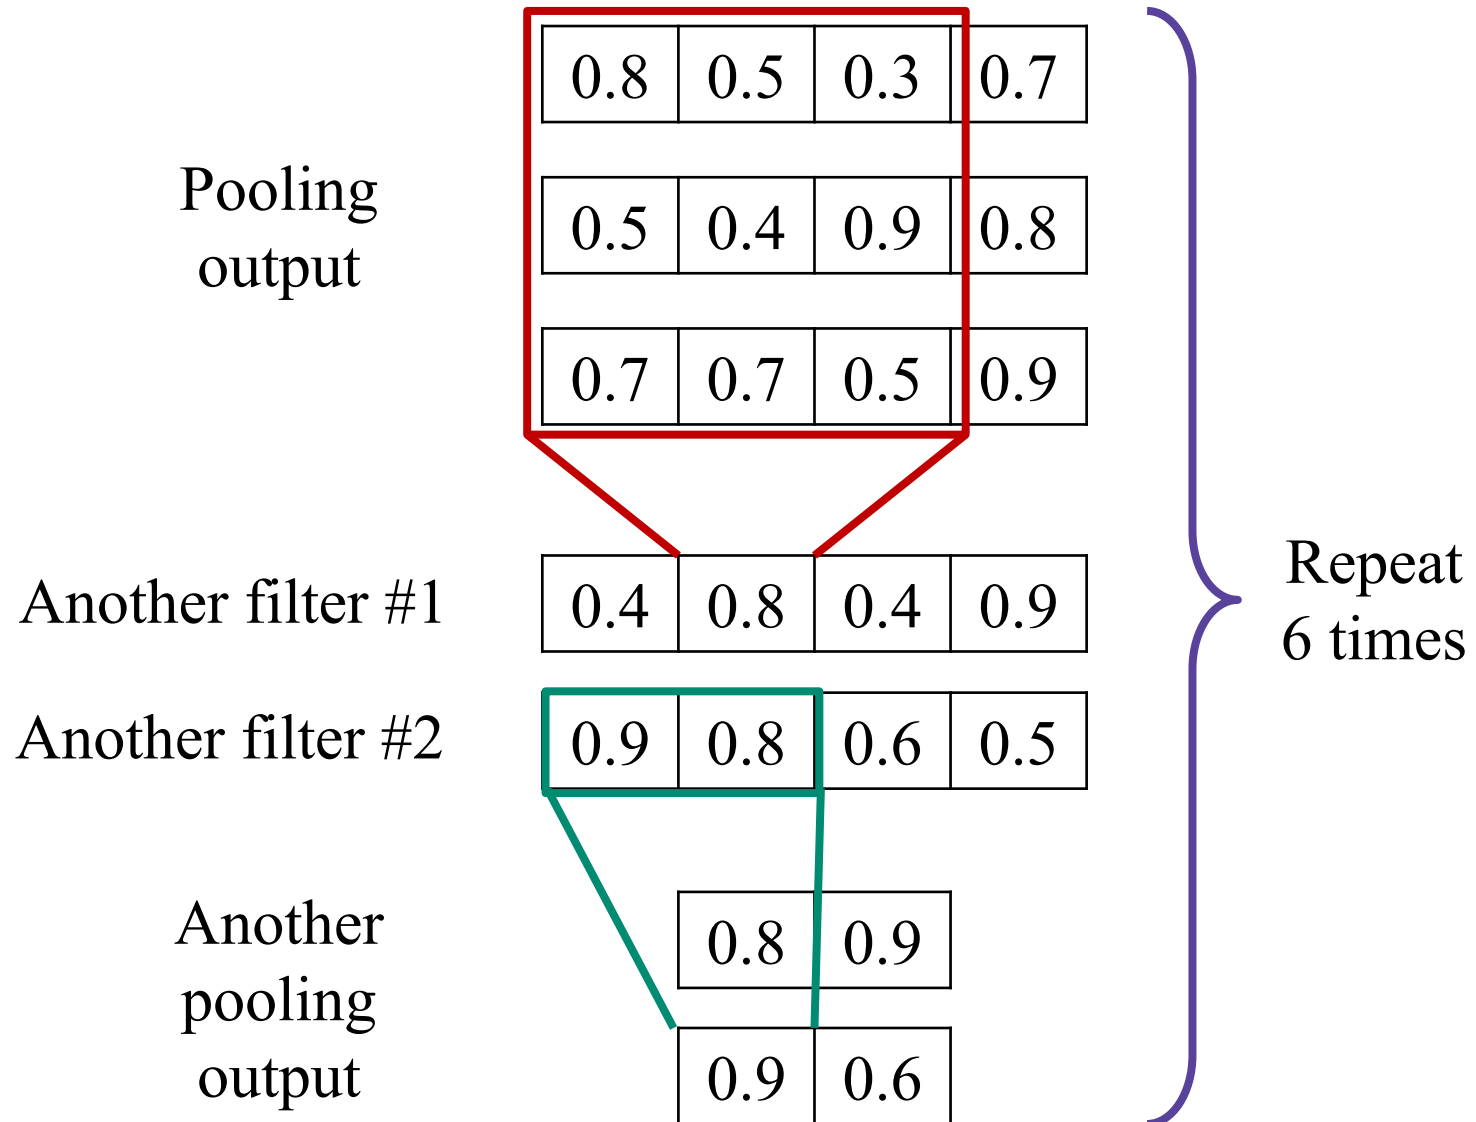
Another filter #2

0.9	0.8	0.6	0.5
-----	-----	-----	-----

Another  
pooling  
output

0.8	0.9
0.9	0.6

# Repeat 1D convolution + pooling





# Final architecture

- Let's take only first **1014** characters of text
- Apply 1D convolution + max pooling **6** times
  - Kernels widths: 7, 7, 3, 3, 3, 3
  - Filters at each step: 1024
- After that we have a **1024** × **34** matrix of features
- Apply MLP for your task

# Experimental datasets

Categorization or sentiment analysis

	Dataset	Classes	Train Samples
<b>Smaller</b>	AG's News	4	120,000
	Sogou News	5	450,000
	DBPedia	14	560,000
	Yelp Review Polarity	2	560,000
<b>Bigger</b>	Yelp Review Full	5	650,000
	Yahoo! Answers	10	1,400,000
	Amazon Review Full	5	3,000,000
	Amazon Review Polarity	2	3,600,000

# Experimental results

Errors on test set for classical models:

Model	AG	Sogou	DBP.	Yelp P.	Yelp F.	Yah. A.	Amz. F.	Amz. P.
BoW	11.19	7.15	3.39	7.76	42.01	31.11	45.36	9.60
BoW TFIDF	10.36	6.55	2.63	6.34	40.14	28.96	44.74	9.00
ngrams	7.96	2.92	1.37	<b>4.36</b>	43.74	31.53	45.73	7.98
ngrams TFIDF	<b>7.64</b>	<b>2.81</b>	<b>1.31</b>	4.56	45.20	31.49	47.56	8.46

Errors on test set for deep models:

LSTM	13.94	4.82	1.45	5.26	41.83	29.16	40.57	6.10
Sm. Full Conv.	11.59	8.95	1.89	5.67	38.82	30.01	40.88	5.78
Lg. Full Conv. Th.	9.51	-	1.55	4.88	38.04	29.58	40.54	5.51
Sm. Full Conv. Th.	10.89	-	1.69	5.42	<b>37.95</b>	29.90	40.53	5.66
Lg. Conv.	12.82	4.88	1.73	5.89	39.62	29.55	41.31	5.51
Sm. Conv.	15.65	8.65	1.98	6.53	40.84	29.84	40.53	5.50
Lg. Conv. Th.	13.39	-	1.60	5.82	39.30	<b>28.80</b>	40.45	<b>4.93</b>
Sm. Conv. Th.	14.80	-	1.85	6.49	40.16	29.84	<b>40.43</b>	5.67

Deep models work better for large datasets!

# Summary

- You can use convolutional networks on top of characters (called learning from scratch)
- It works best for large datasets where it beats classical approaches (like BOW)
- Sometimes it even beats LSTM that works on word level